Lecture 2: Supervised vs. unsupervised learning, bias-variance tradeoff

Reading: Chapter 2

STATS 202: Data mining and analysis

Jonathan Taylor, 9/26
Slide credits: Sergio Bacallado

Supervised vs. unsupervised learning

In **unsupervised learning** we start with a data matrix:

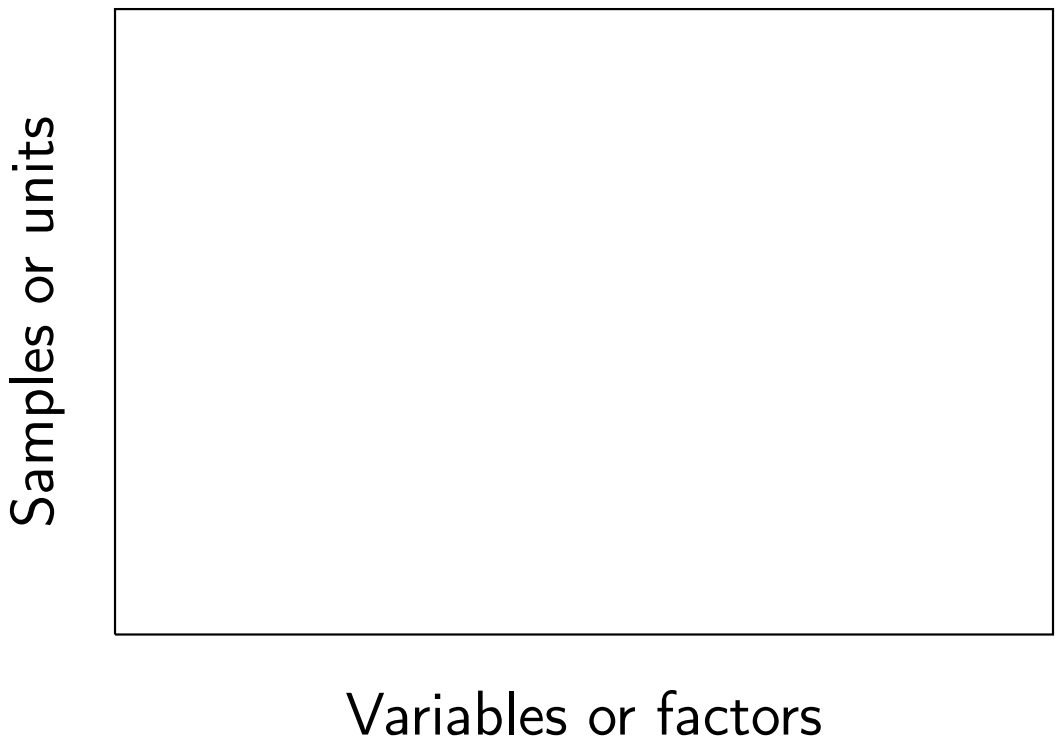

Supervised vs. unsupervised learning

In **unsupervised learning** we start with a data matrix:

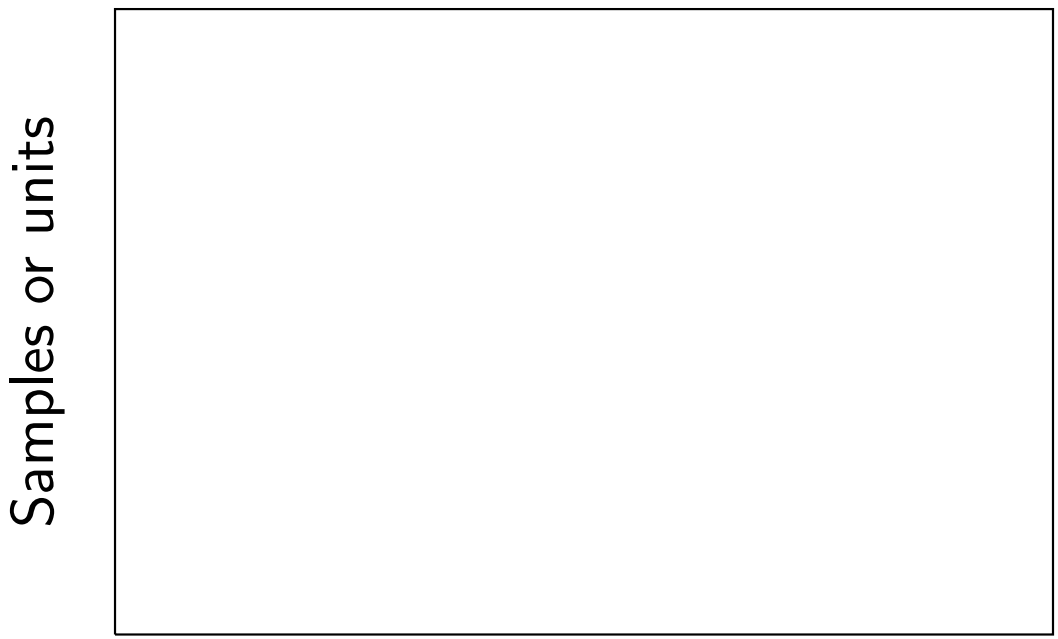

Quantitative, eg. weight, height, number of children, ...

Supervised vs. unsupervised learning

In **unsupervised learning** we start with a data matrix:

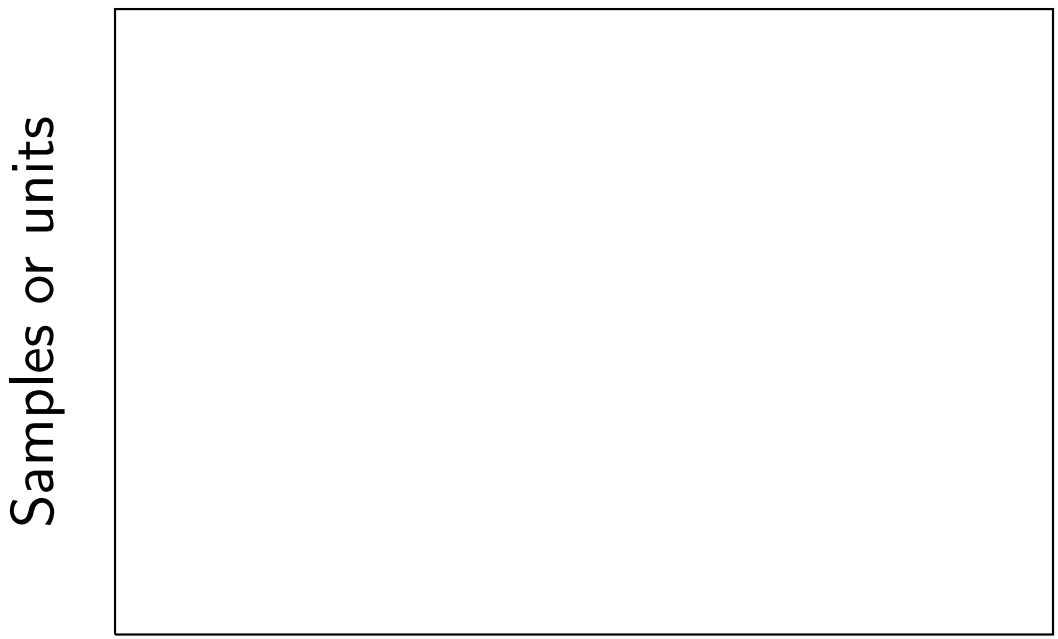

Qualitative, eg. college major, profession, gender, ...

Supervised vs. unsupervised learning

In **unsupervised learning** we start with a data matrix:

Our goal is to:

- Find meaningful relationships between the variables or units.

Supervised vs. unsupervised learning

In **unsupervised learning** we start with a data matrix:

Our goal is to:

- Find meaningful relationships between the variables or units.
- Find low-dimensional representations of the data which make it easy to visualize the variables and units.

Supervised vs. unsupervised learning

In **unsupervised learning** we start with a data matrix:

Our goal is to:

- Find meaningful relationships between the variables or units.
- Find low-dimensional representations of the data which make it easy to visualize the variables and units.
- Find meaningful groupings of the data.

Supervised vs. unsupervised learning

In **unsupervised learning** we start with a data matrix:

Our goal is to:

- Find meaningful relationships between the variables or units. Correlation analysis.
- Find low-dimensional representations of the data which make it easy to visualize the variables and units. PCA, ICA, isomap, locally linear embeddings, etc.
- Find meaningful groupings of the data. Clustering.

Supervised vs. unsupervised learning

In **unsupervised learning** we start with a data matrix:

Our goal is to:

- Find meaningful relationships between the variables or units.
- Find low-dimensional representations of the data which make it easy to visualize the variables and units.
- Find meaningful groupings of the data.

Unsupervised learning is also known in Statistics as **exploratory data analysis**.

Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:

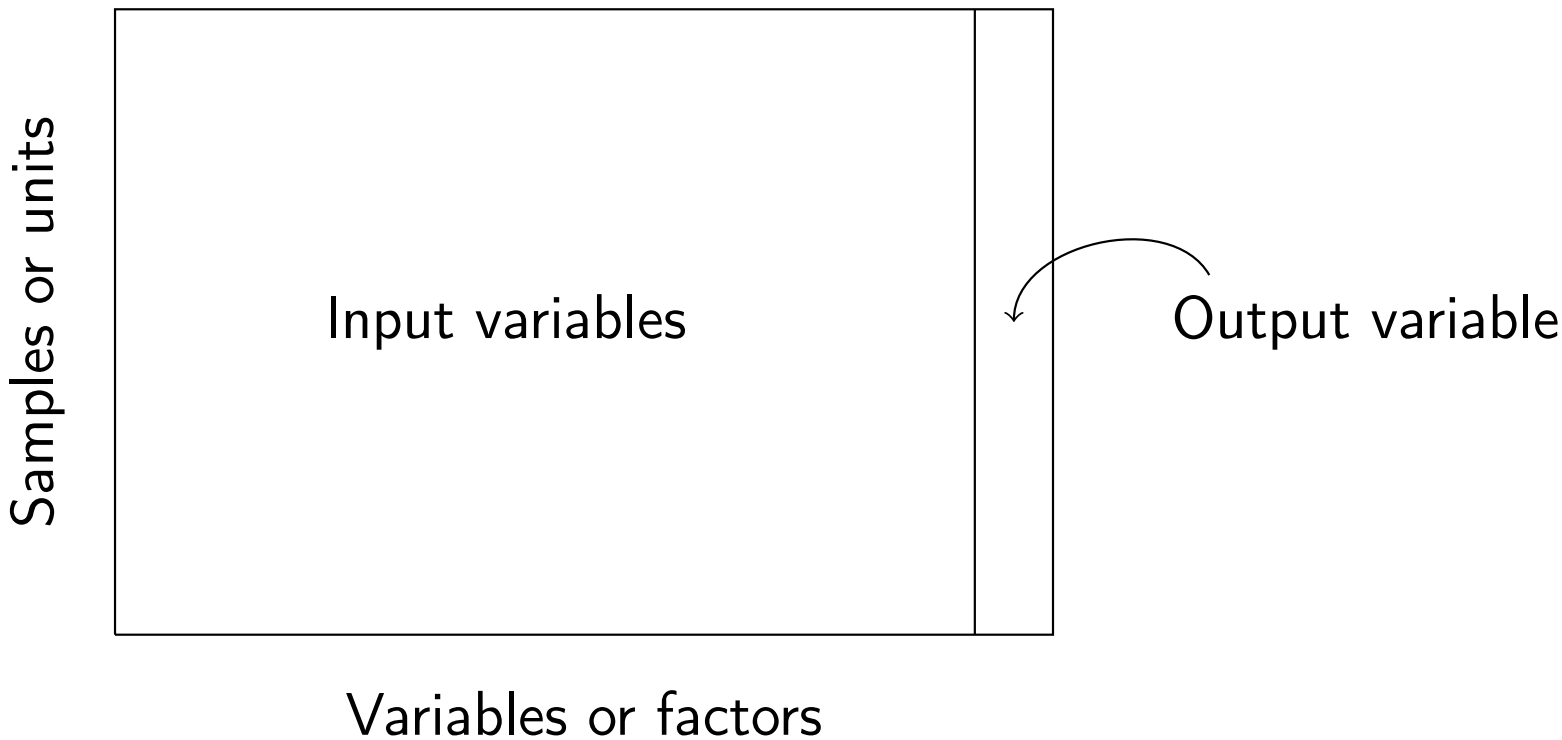

Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:

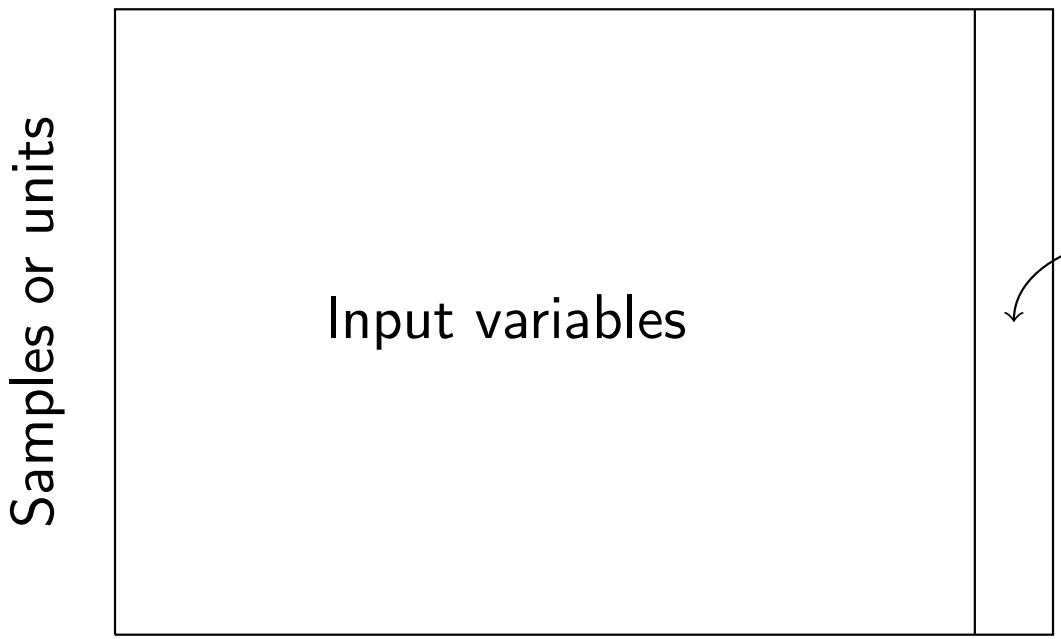

Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:

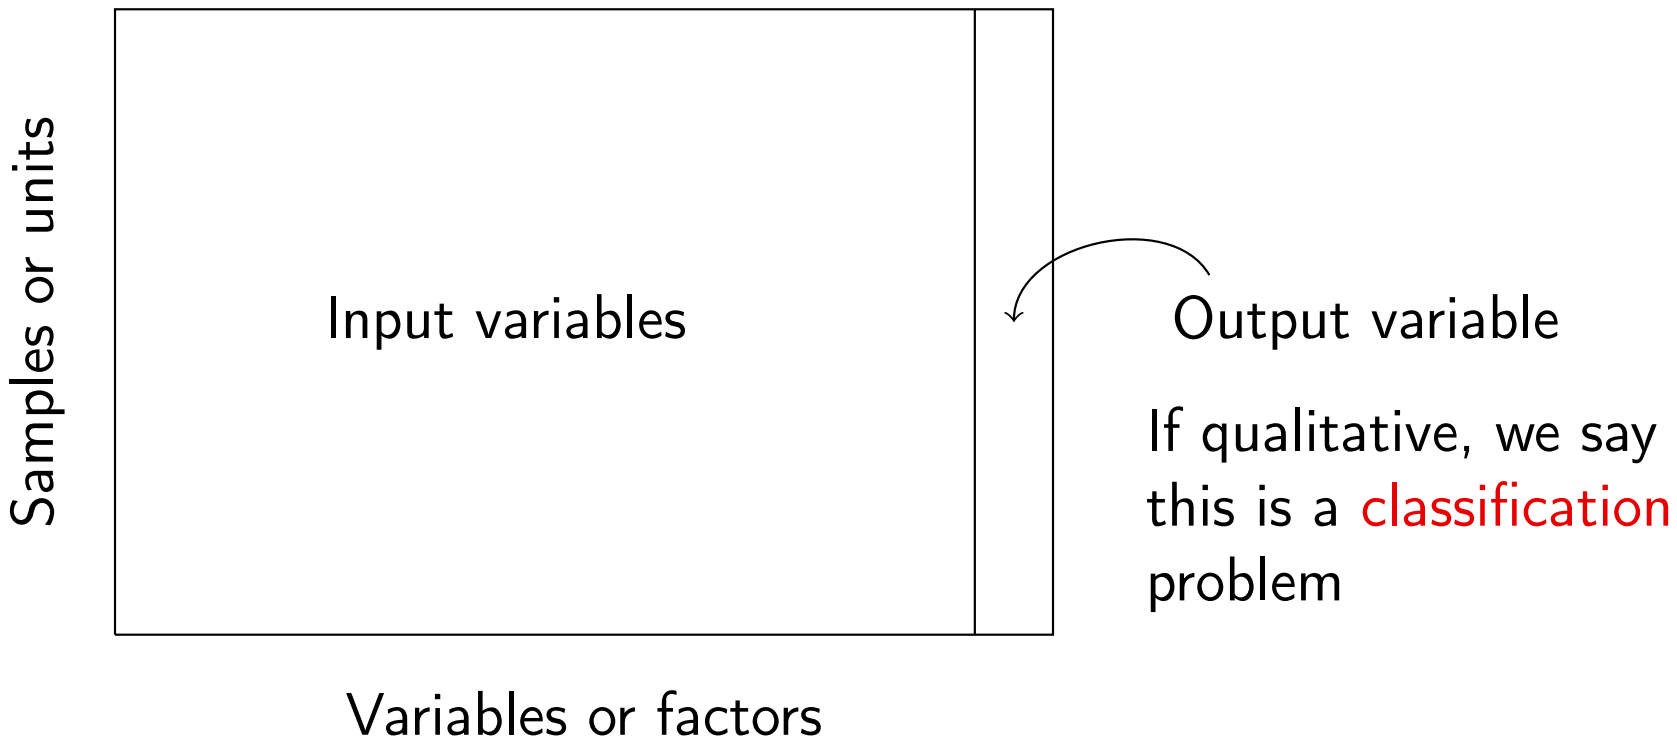

Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:

If X is the vector of inputs for a particular sample. The output variable is modeled by:

$$Y = f(X) + \underbrace{\varepsilon}_{\text{Random error}}$$

Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:

If X is the vector of inputs for a particular sample. The output variable is modeled by:

$$Y = f(X) + \underbrace{\varepsilon}_{\text{Random error}}$$

Our goal is to learn the function f, using a set of training samples.

Typically we assume ϵ independent of X given $f(X)$, i.e. $Y|X \sim N(f(X), \sigma^2)$ though this is not necessary.

Supervised vs. unsupervised learning

$$Y = f(X) + \underbrace{\varepsilon}_{\text{Random error}}$$

Motivations:

- **Prediction:** Useful when the input variable is readily available, but the output variable is not.

  Example: Predict stock prices next month using data from last year.

Supervised vs. unsupervised learning

$$Y = f(X) + \underbrace{\varepsilon}_{\text{Random error}}$$

Motivations:

- **Prediction:** Useful when the input variable is readily available, but the output variable is not.
- **Inference:** A model for f can help us understand the structure of the data — which variables influence the output, and which don't? What is the relationship between each variable and the output, e.g. linear, non-linear?

  Example: What is the influence of genetic variations on the incidence of heart disease.

Parametric and nonparametric methods:

There are (broadly) two kinds of supervised learning method:

Parametric and nonparametric methods:

There are (broadly) two kinds of supervised learning method:

- **Parametric methods:** We assume that f takes a specific form. For example, a linear form:

$$f(X) = X_1\beta_1 + \cdots + X_p\beta_p$$

with parameters $\beta_1, \ldots, \beta_p$. Using the training data, we try to *fit* the parameters.

Parametric and nonparametric methods:

There are (broadly) two kinds of supervised learning method:

- **Parametric methods:** We assume that f takes a specific form. For example, a linear form:

$$f(X) = X_1\beta_1 + \cdots + X_p\beta_p$$

  with parameters $\beta_1, \ldots, \beta_p$. Using the training data, we try to *fit* the parameters.

- **Non-parametric methods:** We don't make any assumptions on the form of f, but we restrict how "wiggly" or "rough" the function can be.

Parametric vs. nonparametric prediction

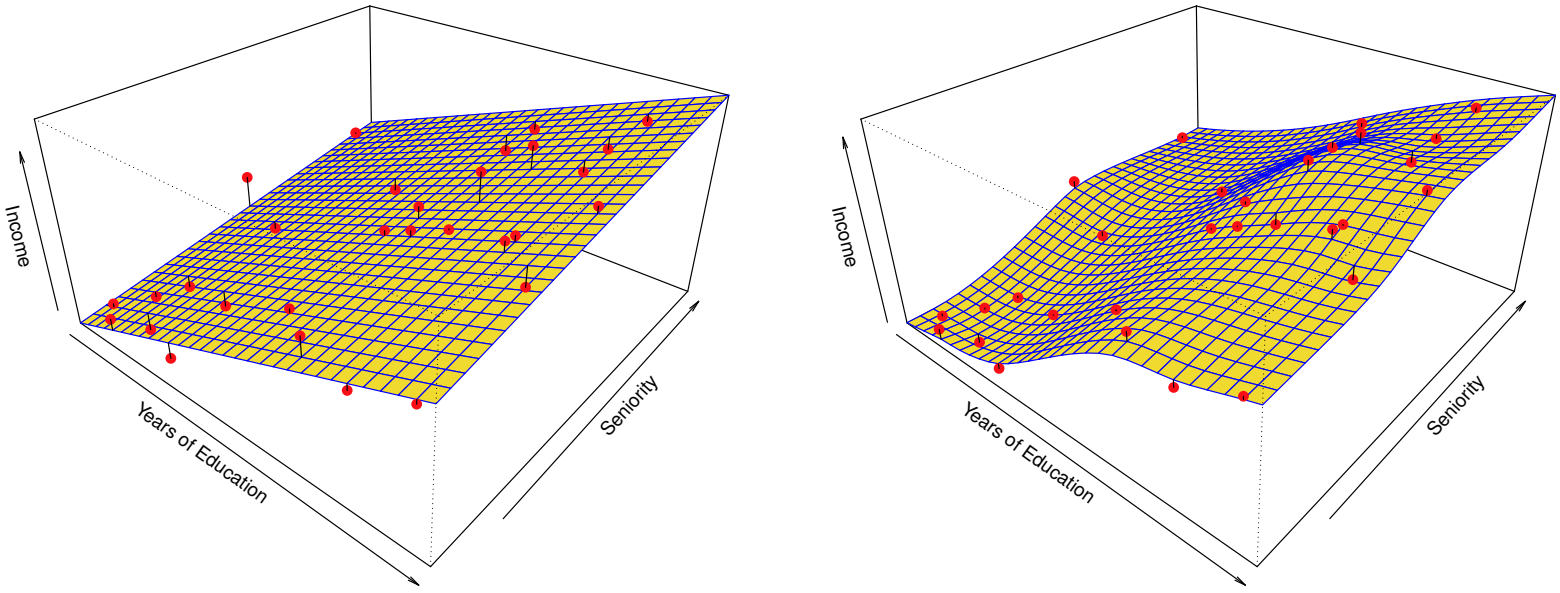

Figures 2.4 and 2.5

Parametric vs. nonparametric prediction

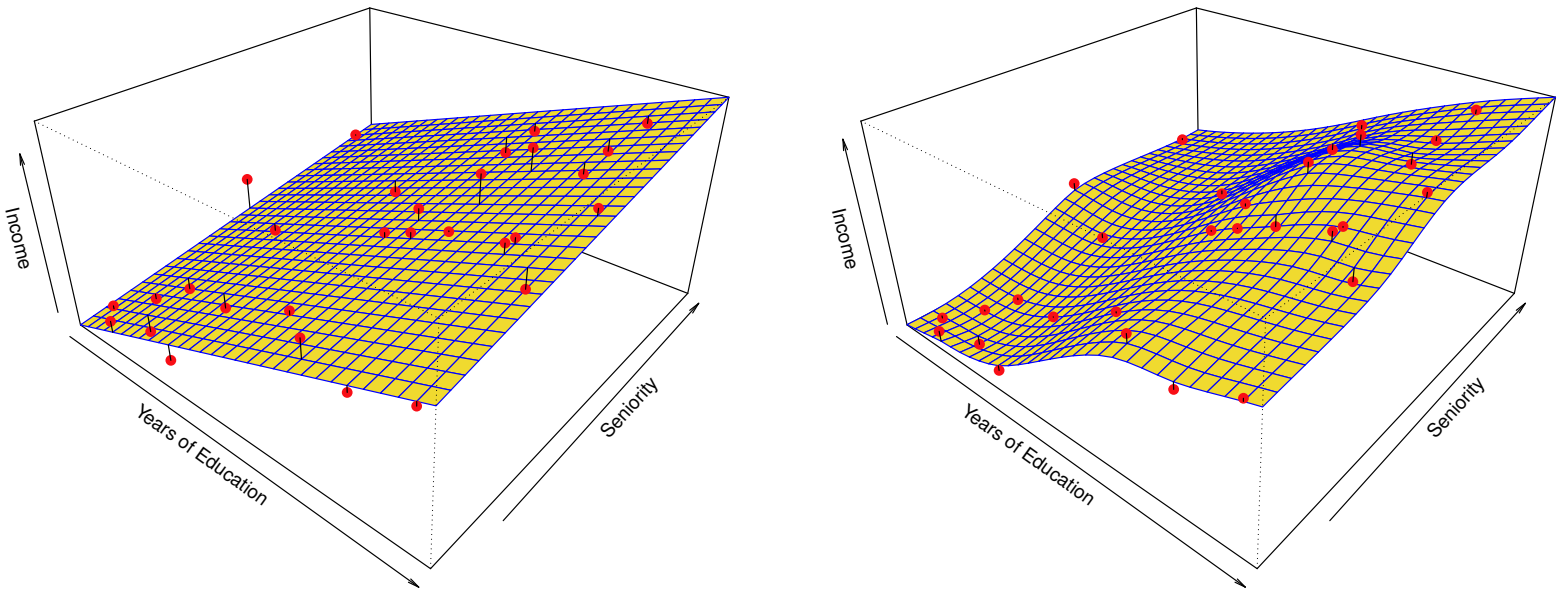

Figures 2.4 and 2.5

Parametric methods have a limit of fit quality. Non-parametric methods keep improving as we add more data to fit.

Parametric vs. nonparametric prediction

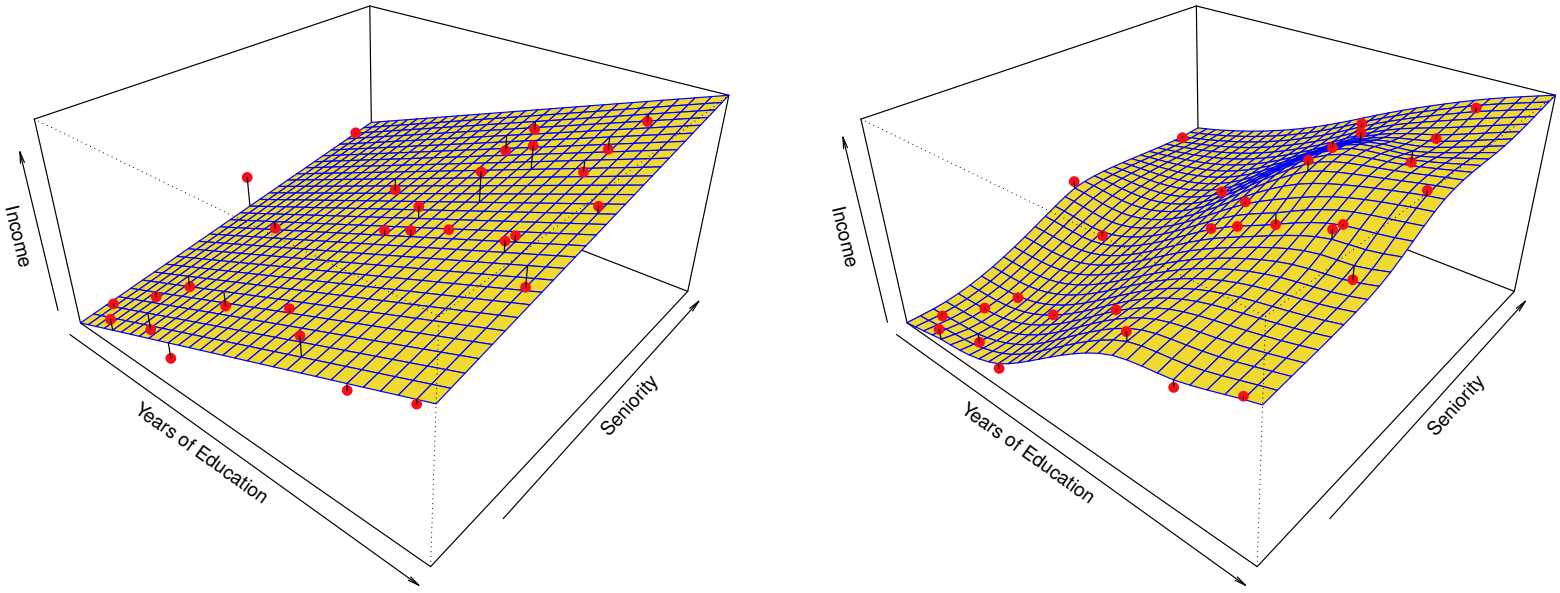

Figures 2.4 and 2.5

Parametric methods have a limit of fit quality. Non-parametric methods keep improving as we add more data to fit.

Parametric methods are often simpler to interpret.

Prediction error

Training data: $(x_1, y_1), (x_2, y_2) \ldots (x_n, y_n)$
Predicted function: $\hat{f}$.

Future data: (x_0, y_0) having some distribution.

Our goal in supervised learning is to minimize the prediction error.

Prediction error

Training data: $(x_1, y_1), (x_2, y_2) \ldots (x_n, y_n)$
Predicted function: $\hat{f}$.

Future data: (x_0, y_0) having some distribution.

Our goal in supervised learning is to minimize the prediction error.
For regression models, this is typically the *Mean Squared Error*:

$$MSE(\hat{f}) = E(y_0 - \hat{f}(x_0))^2.$$

Prediction error

Training data: $(x_1, y_1), (x_2, y_2) \ldots (x_n, y_n)$
Predicted function: $\hat{f}$.

Future data: (x_0, y_0) having some distribution.

Our goal in supervised learning is to minimize the prediction error. For regression models, this is typically the *Mean Squared Error*:

$$MSE(\hat{f}) = E(y_0 - \hat{f}(x_0))^2.$$

Unfortunately, this quantity cannot be computed, because we don't know the joint distribution of (X, Y).

Prediction error

Training data: $(x_1, y_1), (x_2, y_2) \ldots (x_n, y_n)$
Predicted function: $\hat{f}$.

Future data: (x_0, y_0) having some distribution.

Our goal in supervised learning is to minimize the prediction error. For regression models, this is typically the *Mean Squared Error*:

$$MSE(\hat{f}) = E(y_0 - \hat{f}(x_0))^2.$$

Unfortunately, this quantity cannot be computed, because we don't know the joint distribution of (X, Y). We can compute a sample average using the training data; this is known as the training MSE:

$$MSE_{\text{training}}(\hat{f}) = \frac{1}{n} \sum_{i=1}^{n} (y_i - \hat{f}(x_i))^2.$$

Prediction error

The main challenge of statistical learning is that *a low training MSE does not imply a low MSE.*

Prediction error

The main challenge of statistical learning is that *a low training MSE does not imply a low MSE.*

If we have test data $\{(x_i', y_i'); i = 1, \ldots, m\}$ which were not used to fit the model, a better measure of quality for $\hat{f}$ is the test MSE:

$$MSE_{\text{test}}(\hat{f}) = \frac{1}{m} \sum_{i=1}^{m} (y_i' - \hat{f}(x_i'))^2.$$

Figure 2.9.

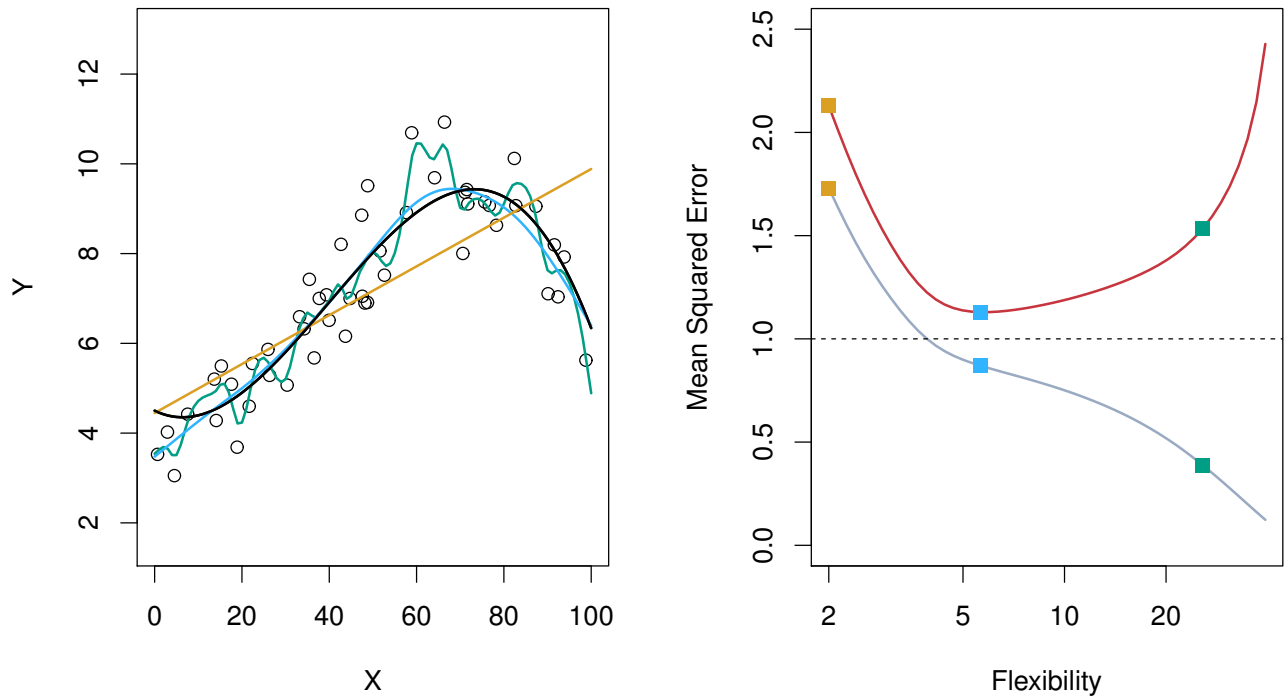

The circles are simulated data from the black curve by adding Gaussian noise.

Figure 2.9.

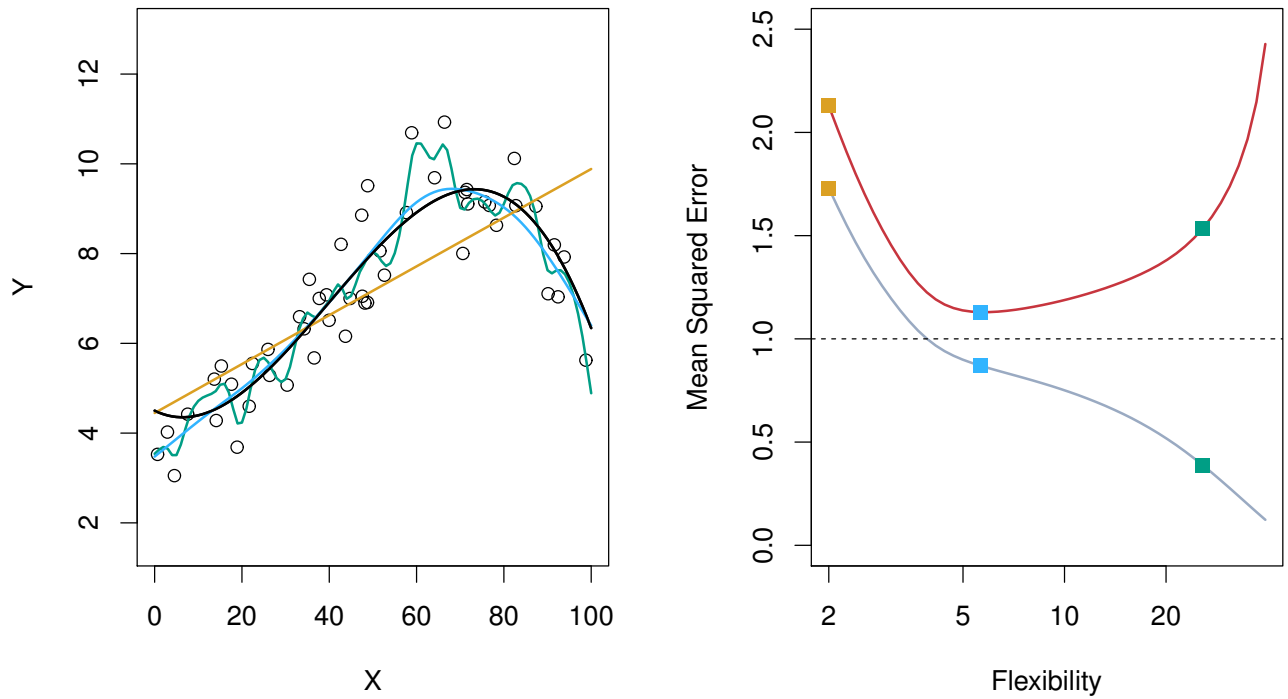

The circles are simulated data from the black curve by adding Gaussian noise. In this artificial example, we *know* what f is.

Figure 2.9.

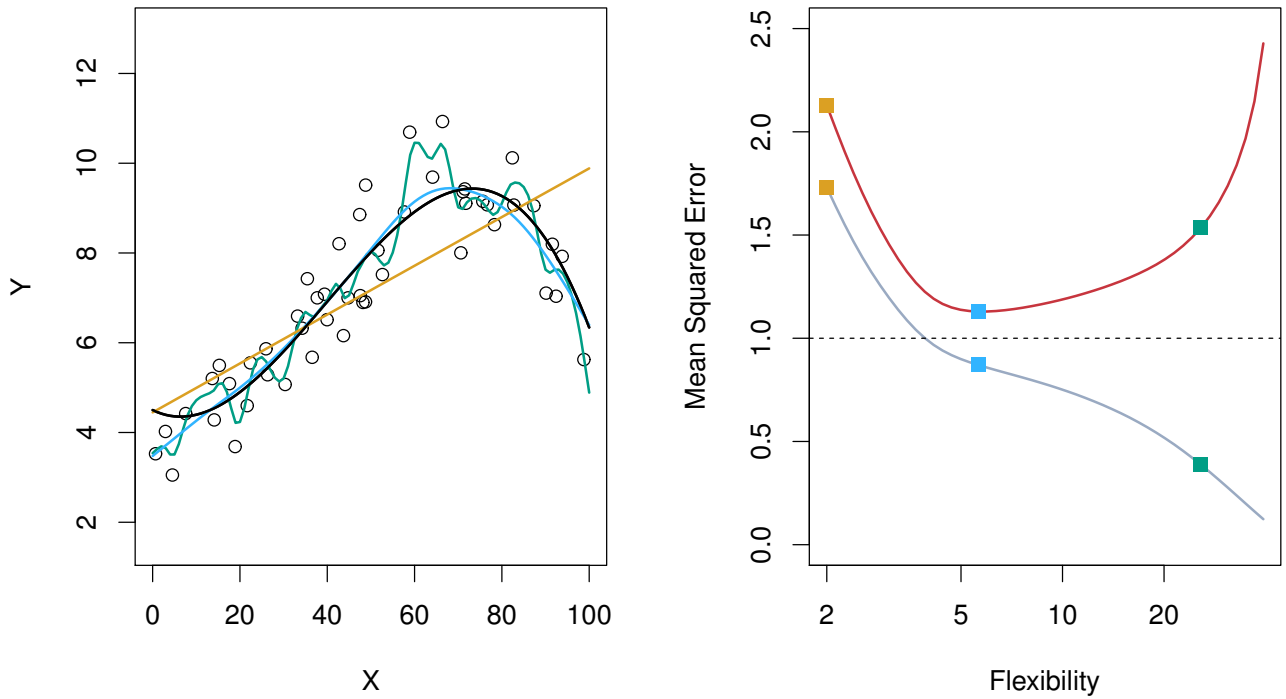

Three estimates $\hat{f}$ are shown:

1. Linear regression.
2. Splines (very smooth).
3. Splines (quite rough).

Figure 2.9.

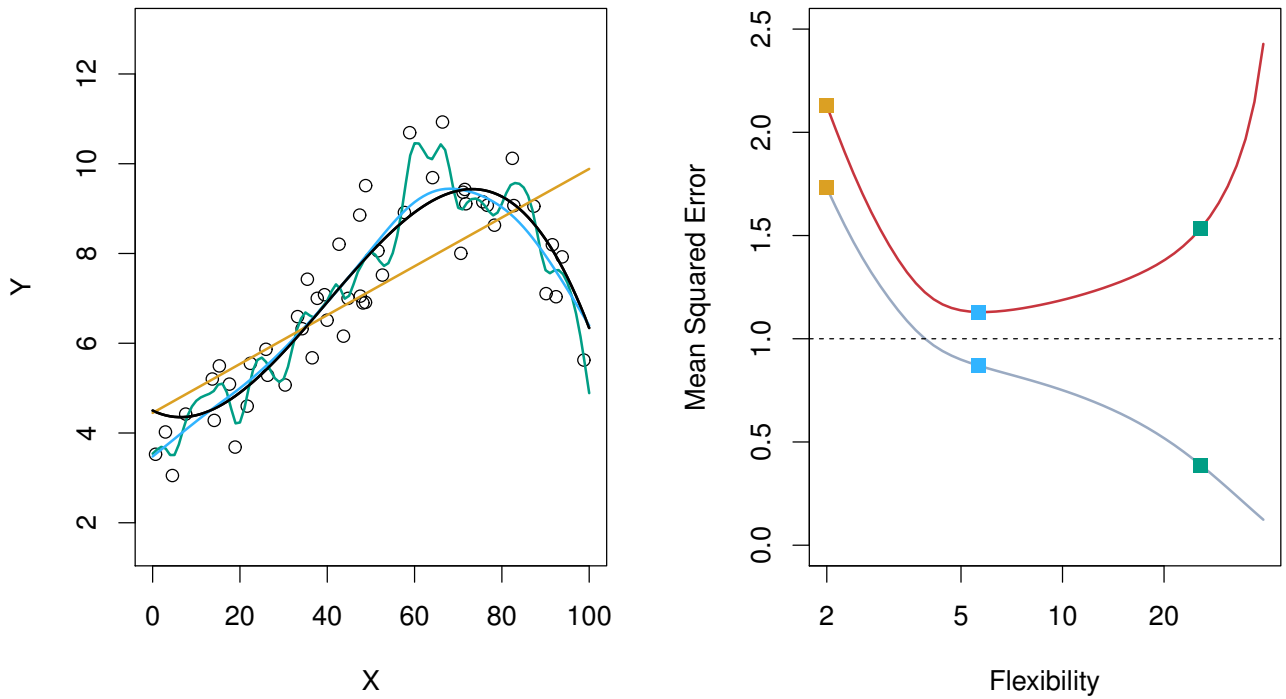

Red line: Test MSE.

Gray line: Training MSE.

Figure 2.10

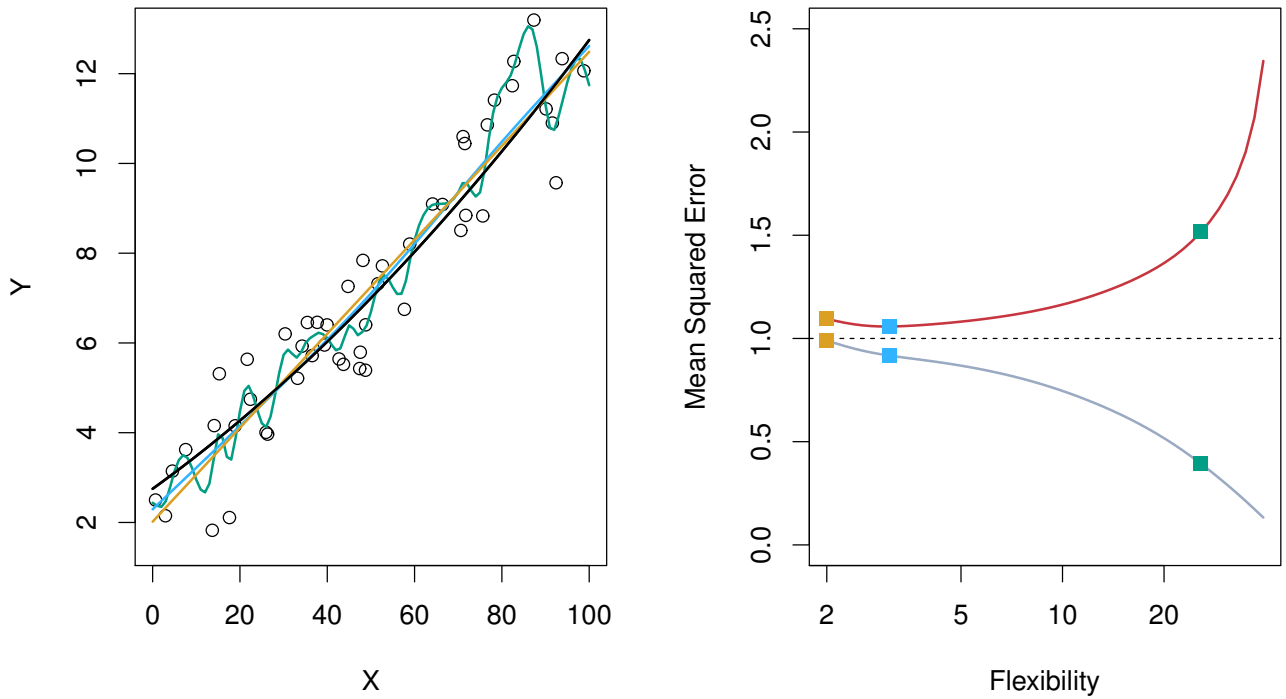

The function f is now almost linear.

Figure 2.11

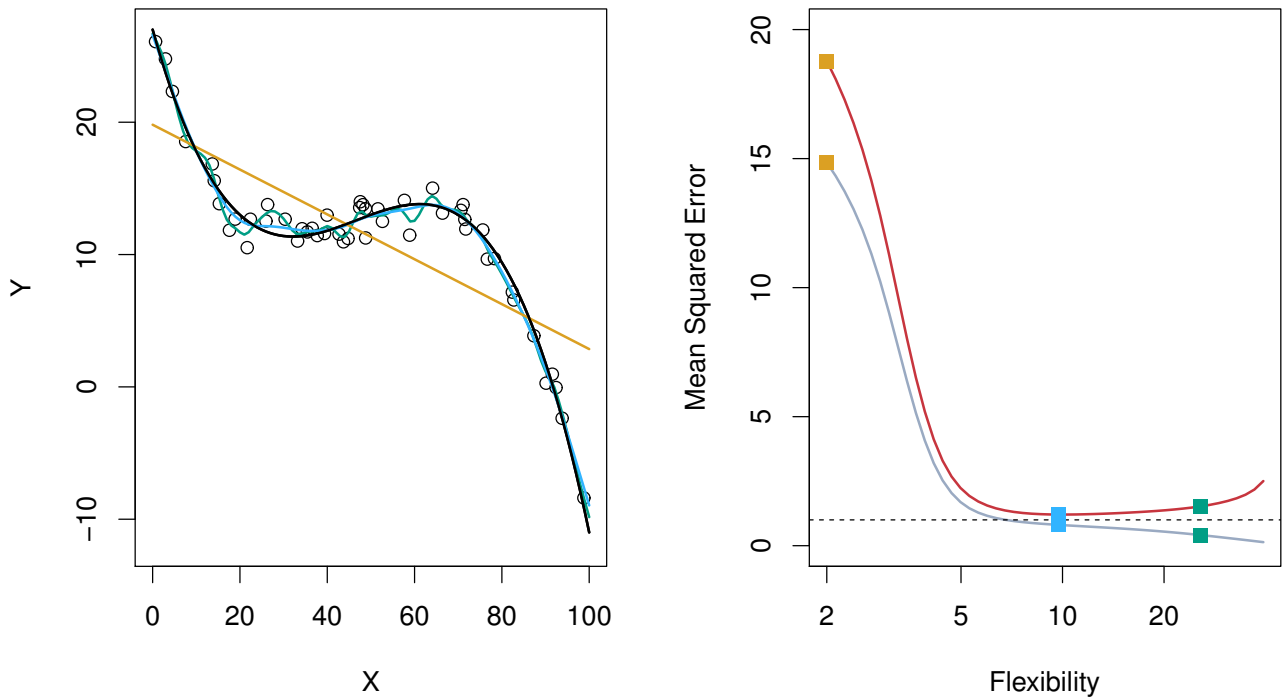

When the noise ε has small variance relative to f, the third method does well.

The bias variance decomposition

Let x_0 be a fixed test point, $y_0 = f(x_0) + \varepsilon_0$, and $\hat{f}$ be estimated from n training samples $(x_1, y_1) \ldots (x_n, y_n)$.

Let E denote the expectation over y_0 and the training outputs $(y_1, \ldots, y_n)$. Then, the Mean Squared Error at x_0 can be decomposed:

$$MSE(x_0) = E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon_0).$$

The bias variance decomposition

Let x_0 be a fixed test point, $y_0 = f(x_0) + \varepsilon_0$, and $\hat{f}$ be estimated from n training samples $(x_1, y_1) \ldots (x_n, y_n)$.

Let E denote the expectation over y_0 and the training outputs $(y_1, \ldots, y_n)$. Then, the Mean Squared Error at x_0 can be decomposed:

$$MSE(x_0) = E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon_0).$$

Irreducible error

The bias variance decomposition

Let x_0 be a fixed test point, $y_0 = f(x_0) + \varepsilon_0$, and $\hat{f}$ be estimated from n training samples $(x_1, y_1) \ldots (x_n, y_n)$.

Let E denote the expectation over y_0 and the training outputs $(y_1, \ldots, y_n)$. Then, the Mean Squared Error at x_0 can be decomposed:

$$MSE(x_0) = E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon_0).$$

The variance of the estimate of Y: $E[\hat{f}(x_0) - E(\hat{f}(x_0))]^2$

This measures how much the estimate of $\hat{f}$ at x_0 changes when we sample new training data.

The bias variance decomposition

Let x_0 be a fixed test point, $y_0 = f(x_0) + \varepsilon_0$, and $\hat{f}$ be estimated from n training samples $(x_1, y_1) \ldots (x_n, y_n)$.

Let E denote the expectation over y_0 and the training outputs $(y_1, \ldots, y_n)$. Then, the Mean Squared Error at x_0 can be decomposed:

$$MSE(x_0) = E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon_0).$$

The squared bias of the estimate of Y: $[E(\hat{f}(x_0)) - f(x_0)]^2$

This measures the deviation of the average prediction $\hat{f}(x_0)$ from the truth $f(x_0)$.

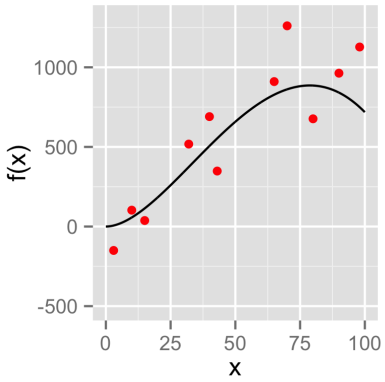
1000
500
0
-500
f(x)
0
25
50
75
100
x

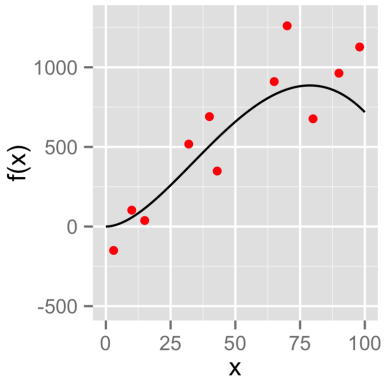
1000
500
0
-500
f(x)
0
25
50
75
100
x

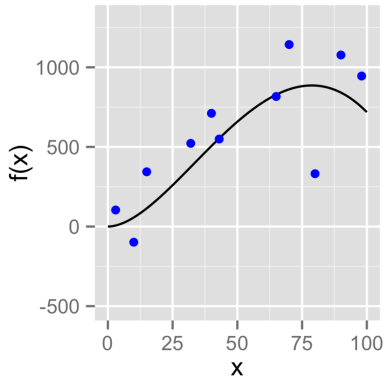
1000
500
0
-500
f(x)
0
25
50
75
100
x

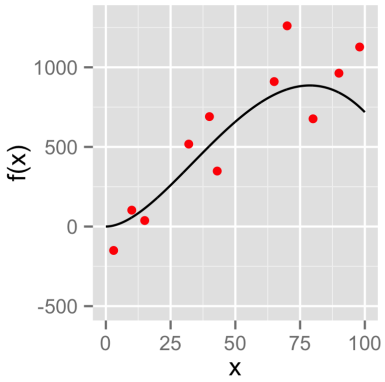
1000
500
0
-500
f(x)
0
25
50
75
100
x

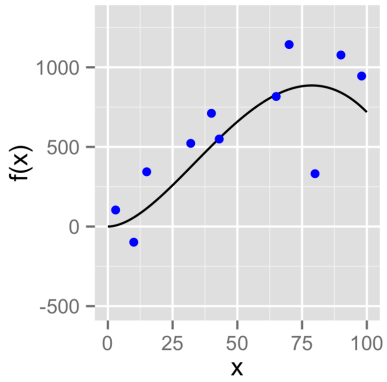
1000
500
0
-500
f(x)
0
25
50
75
100
x

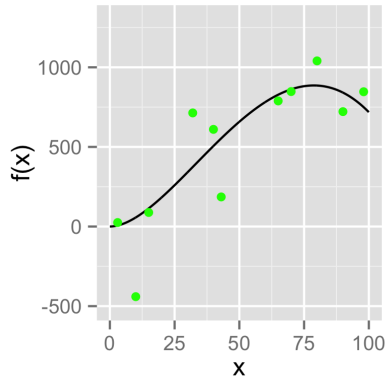
1000
500
0
-500
f(x)
0
25
50
75
100
x

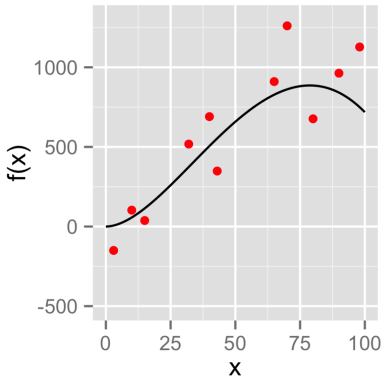
1000
500
0
-500
f(x)
0
25
50
75
100
x

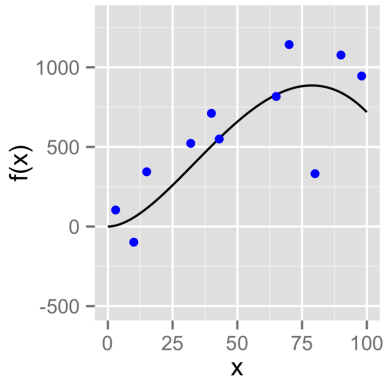
1000
500
0
-500
f(x)
0
25
50
75
100
x

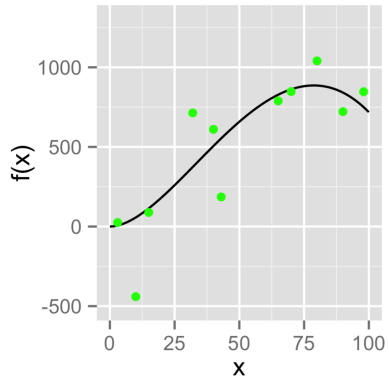
1000
500
0
-500
f(x)
0
25
50
75
100
x

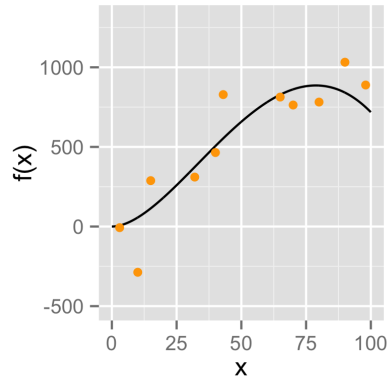
1000
500
0
-500
f(x)
0
25
50
75
100
x

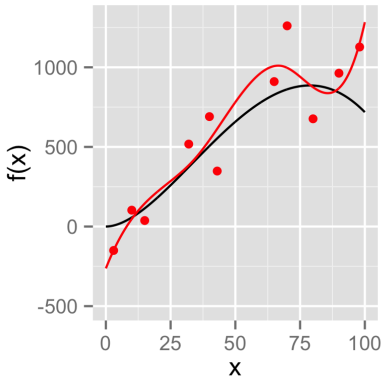
1000
500
0
-500
f(x)
0
25
50
75
100
x

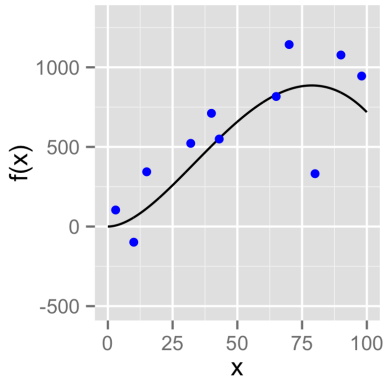
1000
500
0
-500
f(x)
0
25
50
75
100
x

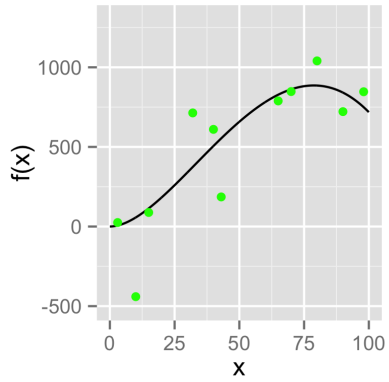
1000
500
0
-500
f(x)
0
25
50
75
100
x

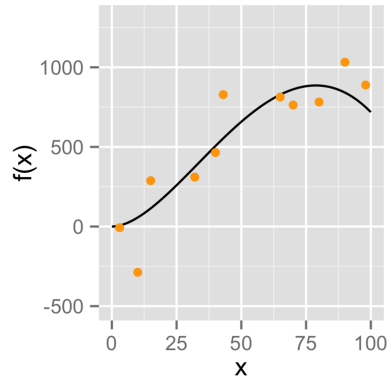
1000
500
0
-500
f(x)
0
25
50
75
100
x

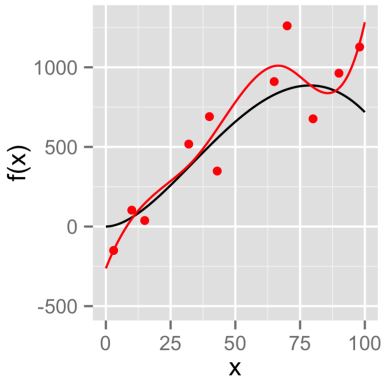
f(x)
1000
500
0
-500
0
25
50
75
100
x

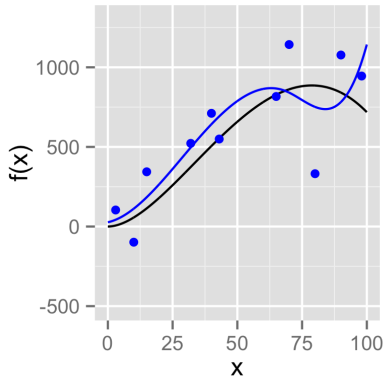
f(x)
1000
500
0
-500
0
25
50
75
100
x

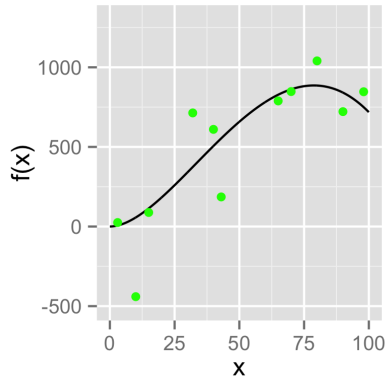
f(x)
1000
500
0
-500
0
25
50
75
100
x

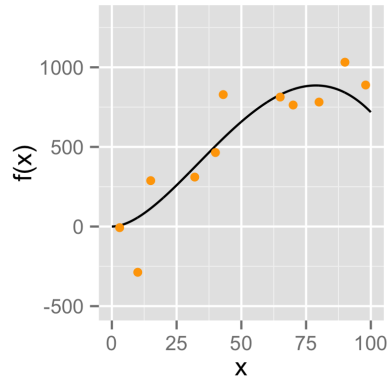
f(x)
1000
500
0
-500
0
25
50
75
100
x

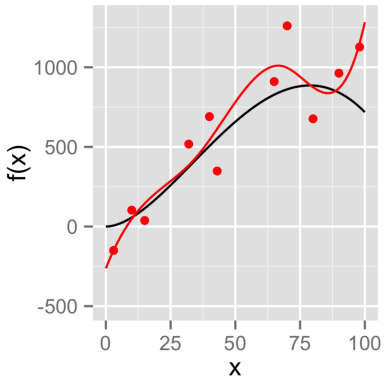
f(x)
1000
500
0
-500
0
25
50
75
100
x

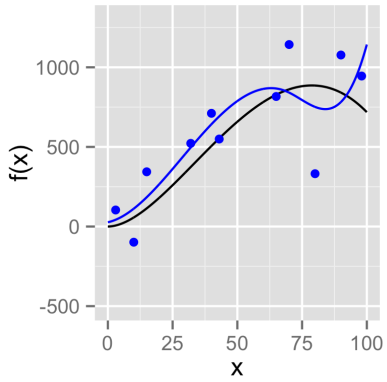
f(x)
1000
500
0
-500
0
25
50
75
100
x

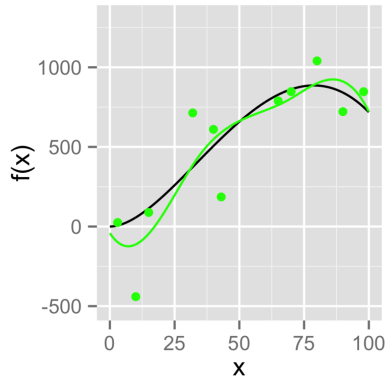
f(x)
1000
500
0
-500
0
25
50
75
100
x

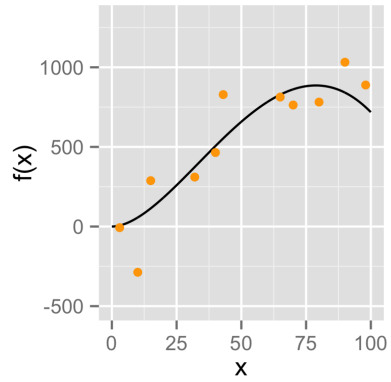
f(x)
1000
500
0
-500
0
25
50
75
100
x

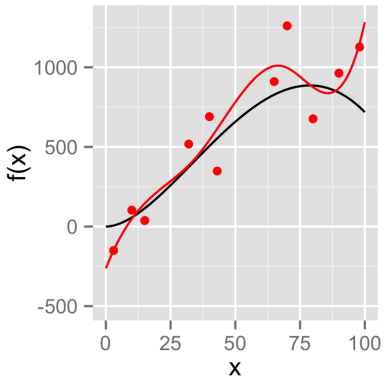
f(x)
1000
500
0
-500
0
25
50
75
100
x

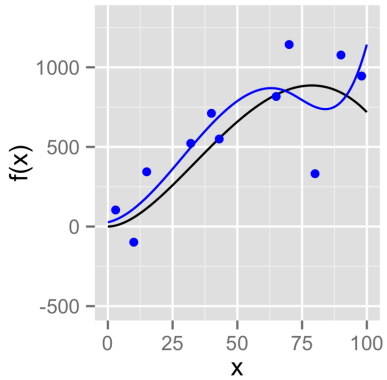
f(x)
1000
500
0
-500
0
25
50
75
100
x

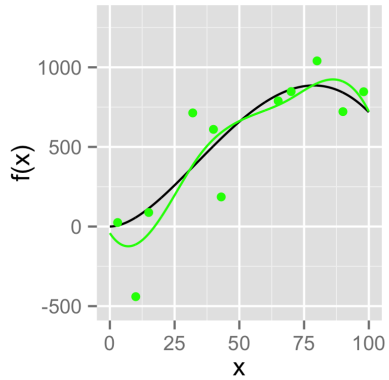
f(x)
1000
500
0
-500
0
25
50
75
100
x

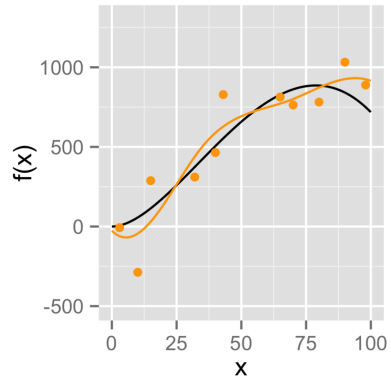
f(x)
1000
500
0
-500
0
25
50
75
100
x

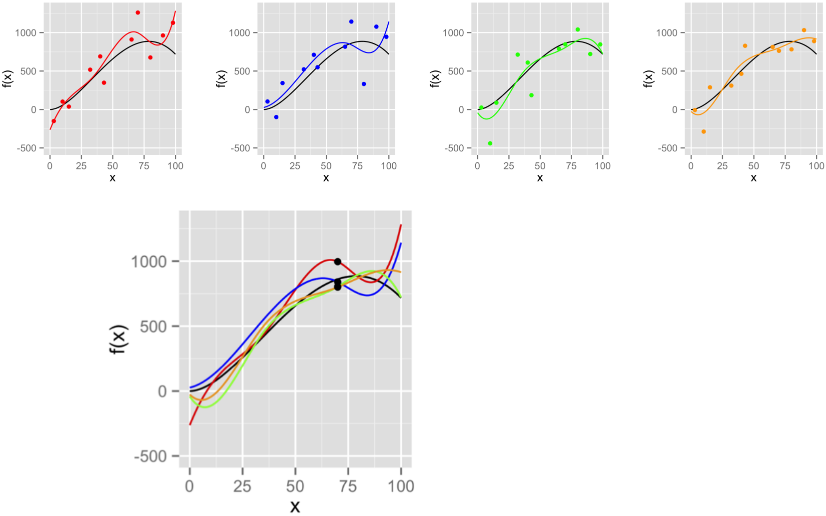
f(x)
x
1000
500
0
-500
0
25
50
75
100

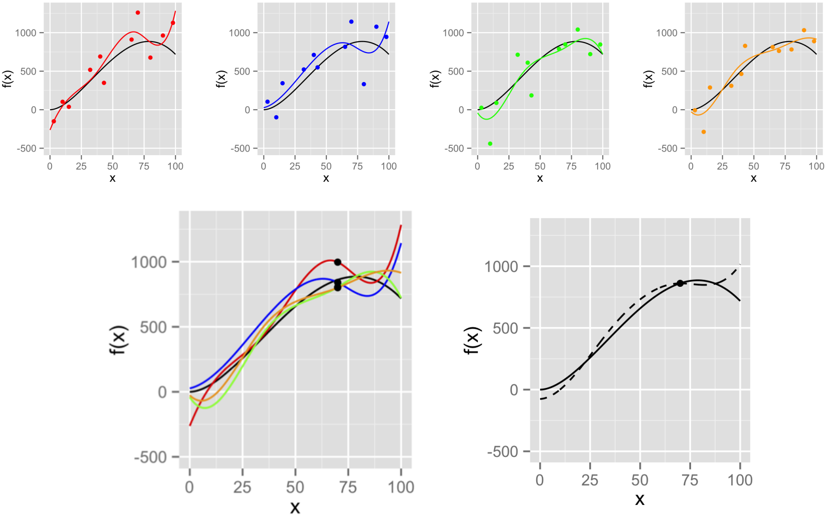

See lab2 to generate these examples.

Implications of bias variance decomposition

$$MSE(x_0) = E(y_0 - \hat{f}(x_0))^2 = \mathsf{Var}(\hat{f}(x_0)) + [\mathsf{Bias}(\hat{f}(x_0))]^2 + \mathsf{Var}(\varepsilon).$$

- The MSE is always positive.

Implications of bias variance decomposition

$$MSE(x_0) = E(y_0 - \hat{f}(x_0))^2 = \mathsf{Var}(\hat{f}(x_0)) + [\mathsf{Bias}(\hat{f}(x_0))]^2 + \mathsf{Var}(\varepsilon).$$

- The MSE is always positive.
- Each element on the right hand side is always positive.

Implications of bias variance decomposition

$$MSE(x_0) = E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon).$$

- The MSE is always positive.
- Each element on the right hand side is always positive.
- Therefore, typically when we decrease the bias beyond some point, we increase the variance, and vice-versa.

Implications of bias variance decomposition

$$MSE(x_0) = E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon).$$

- The MSE is always positive.
- Each element on the right hand side is always positive.
- Therefore, typically when we decrease the bias beyond some point, we increase the variance, and vice-versa.

More flexibility $\iff$ Higher variance $\iff$ Lower bias.

Squiggly f, high noise

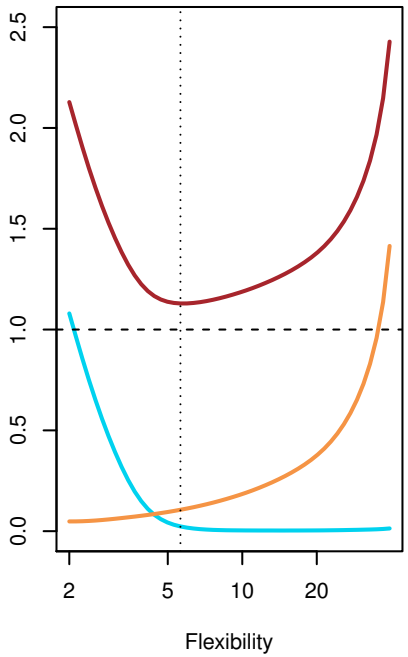

Linear f, high noise

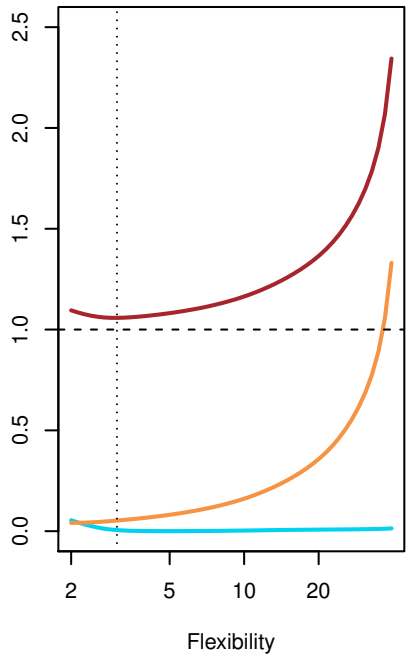

Squiggly f, low noise

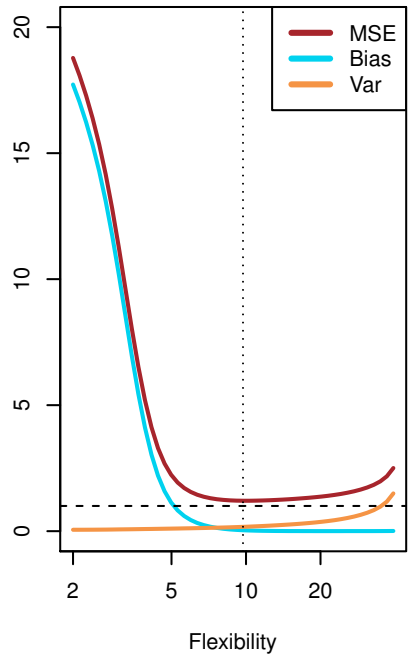

Figure 2.12

Classification problems

In a classification setting, the output takes values in a discrete set.

For example, if we are predicting the brand of a car based on a number of variables, the function f takes values in the set $\{\text{Ford, Toyota, Mercedes-Benz}, \dots\}$.

Classification problems

In a classification setting, the output takes values in a discrete set.

For example, if we are predicting the brand of a car based on a number of variables, the function f takes values in the set $\{\text{Ford, Toyota, Mercedes-Benz}, \dots\}$.

The model:

$$Y = f(X) + \varepsilon$$

becomes insufficient, as X is not necessarily real-valued.

Classification problems

In a classification setting, the output takes values in a discrete set.

For example, if we are predicting the brand of a car based on a number of variables, the function f takes values in the set $\{\text{Ford, Toyota, Mercedes-Benz}, \dots\}$.

The model:

$$\cancel{Y = f(X) + \varepsilon}$$

becomes insufficient, as X is not necessarily real-valued.

Classification problems

In a classification setting, the output takes values in a discrete set.

For example, if we are predicting the brand of a car based on a number of variables, the function f takes values in the set $\{\text{Ford, Toyota, Mercedes-Benz}, \dots\}$.

We will use slightly different notation:

$$P(X, Y) : \text{joint distribution of } (X, Y),$$
$$P(Y \mid X) : \text{conditional distribution of } X \text{ given } Y,$$
$$\hat{y}_i : \text{prediction for } x_i.$$

Loss function for classification

There are many ways to measure the error of a classification prediction. One of the most common is the 0-1 loss:

$$E(\mathbf{1}(y_0 \neq \hat{y}_0))$$

Loss function for classification

There are many ways to measure the error of a classification prediction. One of the most common is the 0-1 loss:

$$E(\mathbf{1}(y_0 \neq \hat{y}_0))$$

Like the MSE, this quantity can be estimated from training and test data by taking a sample average:

$$\frac{1}{n}\sum_{i=1}^{n}\mathbf{1}(y_i \neq \hat{y}_i)$$

Bayes classifier

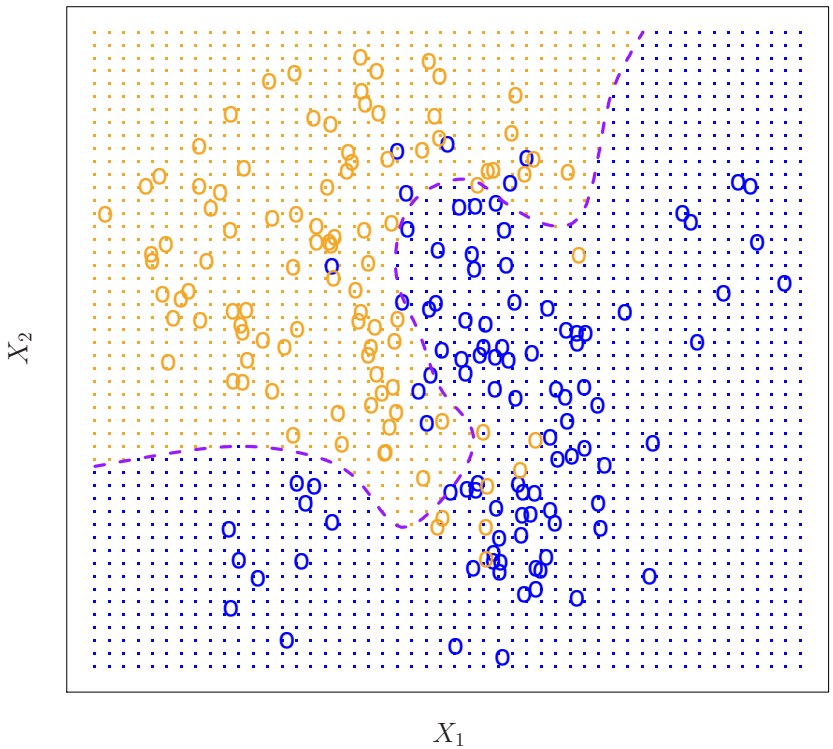

Figure 2.13

In practice, we never know the joint probability P. However, we can assume that it exists.

Bayes classifier

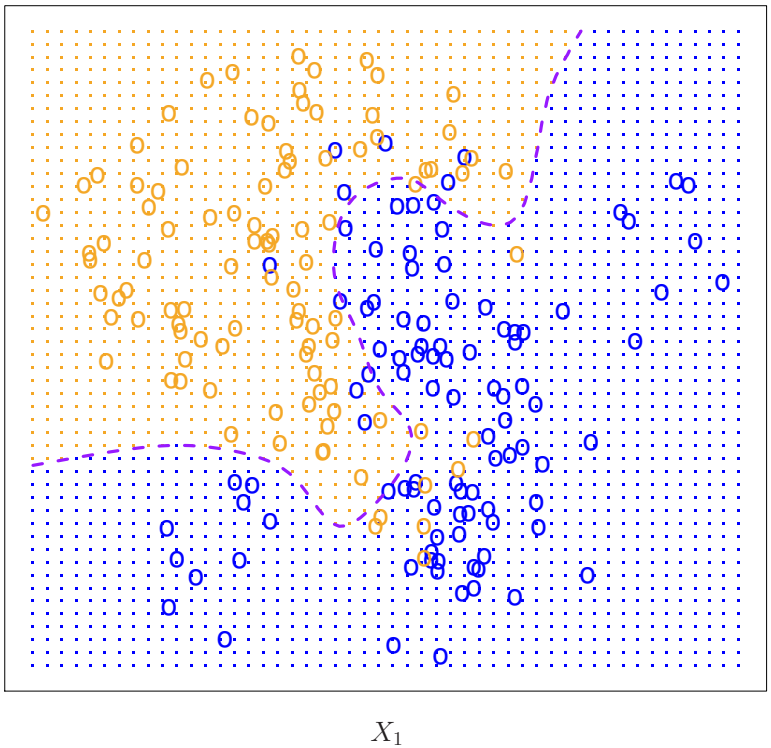

Figure 2.13

In practice, we never know the joint probability P. However, we can assume that it exists.

The Bayes classifier assigns:

$$\hat{y}_i = \text{argmax}_j \;\; P(Y = j \mid X = x_i)$$

It can be shown that this is the best classifier under the 0-1 loss.