

# Regularized Least Squares

Ryan M. Rifkin

Honda Research Institute USA, Inc.  
Human Intention Understanding Group

2007



- Data points  $S = \{(X_1, Y_1), \dots, (X_n, Y_n)\}$ .
- We let  $X$  simultaneously refer to the set  $\{X_1, \dots, X_n\}$  and to the  $n$  by  $d$  matrix whose  $i$ th row is  $X_i^t$ .

- RKHS  $\mathcal{H}$  with a positive semidefinite *kernel function*  $k$ :

linear:  $k(X_i, X_j) = X_i^t X_j$

polynomial:  $k(X_i, X_j) = (X_i^t X_j + 1)^d$

gaussian:  $k(X_i, X_j) = \exp\left(-\frac{\|X_i - X_j\|^2}{\sigma^2}\right)$

- Define the kernel matrix  $K$  to satisfy  $K_{ij} = k(X_i, X_j)$ .
- Abusing notation, allow  $k$  to take and produce sets:
  - $k(X, X) = K$
  - Given an arbitrary point  $X_*$ ,  $k(X, X_*)$  is a column vector whose  $i$ th entry is  $k(X_i, X_*)$ .
- The linear kernel has special properties, which we discuss in detail later.

# The RLS Setup

- Goal: Find the function  $f \in \mathcal{H}$  that minimizes the weighted sum of the *total* square loss and the RKHS norm

$$\min_{f \in \mathcal{H}} \frac{1}{2} \sum_{i=1}^n (f(X_i) - Y_i)^2 + \frac{\lambda}{2} \|f\|_K^2. \quad (1)$$

- Note that in this formulation, we are minimizing the total instead of the average loss. We avoid mucking around with the factor of  $1/n$ , which can be folded into  $\lambda$ .
- This loss function “makes sense” for regression. We can also use it for binary classification, where it “makes no sense” but works great.

# Applying the Representer

- The representer theorem guarantees that the solution to (1) can be written as

$$f(\cdot) = \sum_{i=1}^n c_i k(X_i, \cdot),$$

for some  $c \in \mathbb{R}^n$ .

- We can therefore rewrite (1) as

$$\min_{c \in \mathbb{R}^n} \frac{1}{2} \|Y - Kc\|_2^2 + \frac{\lambda}{2} \|f\|_K^2.$$

# Applying the Representer Theorem, II

Consider a function of the form:

$$f(\cdot) = \sum_{i=1}^n c_i k(X_i, \cdot),$$

For such a function,

$$\begin{aligned} \|f\|_K^2 &= \langle f, f \rangle_K \\ &= \left\langle \sum_{i=1}^n c_i k(X_i, \cdot), \sum_{j=1}^n c_j k(X_j, \cdot) \right\rangle_K \\ &= \sum_{i=1}^n \sum_{j=1}^n c_i c_j \langle k(X_i, \cdot), k(X_j, \cdot) \rangle_K \\ &= \sum_{i=1}^n \sum_{j=1}^n c_i c_j k(X_i, X_j) \\ &= c^t K c \end{aligned}$$

# The RLS Solution



$$\frac{1}{2} \|Y - Kc\|_2^2 + \frac{\lambda}{2} c^t Kc$$

is clearly convex in  $c$  (why?), so we can find its minimum by setting the gradient w.r.t  $c$  to 0:

$$\begin{aligned} -K(Y - Kc) + \lambda Kc &= 0 \\ (K + \lambda I)c &= Y \\ c &= (K + \lambda I)^{-1} Y \end{aligned}$$

- We find  $c$  by solving a system of linear equations.

# The RLS Solution, Comments

- The solution exists and is unique (for  $\lambda > 0$ ).
- Define  $G(\lambda) = K + \lambda I$ . (Often  $\lambda$  is clear from context and we write  $G$ .)
- The prediction at a new test point  $X_*$  is:

$$\begin{aligned} f(X_*) &= \sum c_i k(X_i, X_*) \\ &= k(X, X_*)^t c \\ &= Y^t G^{-1} k(X, X_*). \end{aligned}$$

- The use of  $G^{-1}$  (or other inverses) is formal only. We do *not* recommend taking matrix inverses.



# Solving RLS, Parameters Fixed.

- Situation: All hyperparameters fixed
- We just need to solve a single linear system

$$(K + \lambda I)c = y.$$

- The matrix  $K + \lambda I$  is **symmetric positive definite**, so the appropriate algorithm is Cholesky factorization.
- In Matlab, the “slash” operator seems to be using Cholesky, so you can just write  $c = (K + \lambda I) \backslash Y$ , but to be safe, (or in octave), I suggest  $R = \text{chol}(K + \lambda I)$ ;  $c = (R \backslash (R' \backslash Y))$  ; .

# Solving RLS, Varying $\lambda$

- Situation: We don't know what  $\lambda$  to use, all other hyperparameters fixed.
- Form the eigendecomposition  $K = Q\Lambda Q^t$ , where  $\Lambda$  is diagonal with  $\Lambda_{jj} \geq 0$  and  $QQ^t = I$ .
- 

$$\begin{aligned} G &= K + \lambda I \\ &= Q\Lambda Q^t + \lambda I \\ &= Q(\Lambda + \lambda I)Q^t, \end{aligned}$$

which implies  $G^{-1} = Q(\Lambda + \lambda I)^{-1}Q^t$ .

# Solving RLS, Varying $\lambda$ , Cont'd

- $O(n^3)$  time to solve one (dense) linear system, *or* to compute the eigendecomposition (constant is maybe 4x worse). Given  $Q$  and  $\Lambda$ , we can find  $c(\lambda)$  in  $O(n^2)$  time:

$$c(\lambda) = Q(\Lambda + \lambda I)^{-1} Q^t Y,$$

noting that  $(\Lambda + \lambda I)$  is diagonal.

- Finding  $c(\lambda)$  for many  $\lambda$ 's is (essentially) free!

- We showed how to find  $c(\lambda)$  quickly as we vary  $\lambda$ .
- But how do we decide if a given  $\lambda$  is “good”?
- Simplest idea: Use the training set error.
- Problem: This frequently overfits.
- Other methods are possible, but today we consider validation.
- Validation means checking our function’s behavior on points other than the training set.

# Types of Validation

- If we have a huge amount of data, we could hold back some percentage of our data (30% is typical), and use this *development* set to choose hyperparameters.
- More common is *k-fold cross-validation*, which means a couple of different things:
  - Divide your data into  $k$  equal sets  $S_1, \dots, S_k$ . For each  $i$ , train on the other  $k - 1$  sets and test on the  $i$ th set.
  - A total of  $k$  times, randomly split your data into a training and test set.
- The limit of (the first kind of)  $k$ -fold validation is *leave-one-out cross-validation*.

# Leave-One-Out Cross-Validation

- For each data point  $x_i$ , build a classifier using the remaining  $n - 1$  data points, and measure the error at  $x_i$ .
- Empirically, this seems to be the method of choice when  $n$  is small.
- Problem: We have to build  $n$  different predictors, on data sets of size  $n - 1$ .
- We will now proceed to show that for RLS, obtaining the LOO error is (essentially) free!

# Leave-One-Out CV: Notation

- Define  $S^i$  to be the data set with the  $i$ th point removed:  
 $S^i = \{(X_1, Y_1), \dots, (X_{i-1}, Y_{i-1}), (X_{i+1}, Y_{i+1}), \dots, (X_n, Y_n)\}.$
- The  $i$ th leave-one-out value is  $f_{S^i}(X_i).$
- The  $i$ th leave-one-out error is  $Y_i - f_{S^i}(X_i).$
- Define  $LOOV$  and  $LOOE$  to be the vectors of leave-one-out values and errors over the training set.
- $\|LOOE\|_2^2$  is considered a good empirical proxy for the error on future points, and we often want to choose parameters by minimizing this quantity.

- Imagine (hallucinate) that we already know  $f_{Si}(X_i)$ .
- Define the vector  $Y^i$  via

$$Y_j^i = \begin{cases} Y_j & j \neq i \\ f_{Si}(X_i) & j = i \end{cases}$$



- Suppose we solve a Tikhonov problem with  $Y^i$  instead of  $Y$  as the labels. Then  $f_{Si}$  is the optimizer:

$$\begin{aligned} & \frac{1}{2} \sum_{j=1}^n (Y_j^i - f(X_j))^2 + \frac{\lambda}{2} \|f\|_K^2 \\ & \geq \frac{1}{2} \sum_{j \neq i} (Y_j^i - f(X_j))^2 + \frac{\lambda}{2} \|f\|_K^2 \\ & \geq \frac{1}{2} \sum_{j \neq i} (Y_j^i - f_{Si}(X_j))^2 + \frac{\lambda}{2} \|f_{Si}\|_K^2 \\ & = \frac{1}{2} \sum_{j=1}^n (Y_j^i - f_{Si}(X_j))^2 + \frac{\lambda}{2} \|f_{Si}\|_K^2. \end{aligned}$$

- Therefore,

$$\begin{aligned}c^i &= G^{-1} Y^i \\ f_{Si}(X_i) &= (KG^{-1} Y^i)_i\end{aligned}$$

- This is circular reasoning so far, because we need to know  $f_{Si}(X_i)$  to form  $Y^i$  in the first place.
- However, assuming we have already solved RLS for the whole training set, and we have computed  $f_S(X) = KG^{-1} Y$ , we can do something nice ...

$$\begin{aligned}f_{S^i}(X_i) - f_S(X_i) &= \sum_j (KG^{-1})_{ij} (Y_j^i - Y_j) \\&= (KG^{-1})_{ii} (f_{S^i}(X_i) - Y_i) \\f_{S^i}(X_i) &= \frac{f_S(X_i) - (KG^{-1})_{ii} Y_i}{1 - (KG^{-1})_{ii}} \\&= \frac{(KG^{-1} Y)_i - (KG^{-1})_{ii} Y_i}{1 - (KG^{-1})_{ii}}.\end{aligned}$$

$$LOOV = \frac{KG^{-1}Y - \text{diag}_m(KG^{-1})Y}{\text{diag}_v(I - KG^{-1})},$$

$$LOOE = Y - LOOV$$

$$= Y + \frac{\text{diag}_m(KG^{-1})Y - KG^{-1}Y}{\text{diag}_v(I - KG^{-1})}$$

$$= \frac{\text{diag}_m(I - KG^{-1})Y}{\text{diag}_v(I - KG^{-1})} + \frac{\text{diag}_m(KG^{-1})Y - KG^{-1}Y}{\text{diag}_v(I - KG^{-1})}$$

$$= \frac{Y - KG^{-1}Y}{\text{diag}_v(I - KG^{-1})}.$$

We can simplify our expressions in a way that leads to better computational and numerical properties by noting

$$\begin{aligned}KG^{-1} &= Q\Lambda Q^t Q(\Lambda + \lambda I)^{-1} Q^t \\&= Q\Lambda(\Lambda + \lambda I)^{-1} Q^t \\&= Q(\Lambda + \lambda I - \lambda I)(\Lambda + \lambda I)^{-1} Q^t \\&= I - \lambda G^{-1}.\end{aligned}$$

Substituting into our expression for *LOOE* yields

$$\begin{aligned} LOOE &= \frac{Y - KG^{-1}Y}{\text{diag}_V(I - KG^{-1})} \\ &= \frac{Y - (I - \lambda G^{-1})Y}{\text{diag}_V(I - (I - \lambda G^{-1}))} \\ &= \frac{\lambda G^{-1}Y}{\text{diag}_V(\lambda G^{-1})} \\ &= \frac{G^{-1}Y}{\text{diag}_V(G^{-1})} \\ &= \frac{c}{\text{diag}_V(G^{-1})}. \end{aligned}$$

# The cost of computing $LOOE$

- For RLS, we compute  $LOOE$  via

$$LOOE = \frac{c}{\text{diag}_v(G^{-1})}.$$

- We already showed how to compute  $c(\lambda)$  in  $O(n^2)$  time (given  $K = Q\Lambda Q^t$ ).
- We can also compute a single entry of  $G(\lambda)^{-1}$  in  $O(n)$  time:

$$\begin{aligned} G_{ij}^{-1} &= (Q(\Lambda + \lambda I)^{-1} Q^t)_{ij} \\ &= \sum_{k=1}^n \frac{Q_{ik} Q_{jk}}{\Lambda_{kk} + \lambda}, \end{aligned}$$

and therefore we can compute  $\text{diag}(G^{-1})$ , and compute  $LOOE$ , in  $O(n^2)$  time.

# Summary So Far

- If we can (directly) solve one RLS problem on our data, we can find a good value of  $\lambda$  using LOO optimization at essentially the same cost.
- When can we solve one RLS problem?
- We need to form  $K$ , which takes  $O(n^2 d)$  time and  $O(n^2)$  memory. We need to perform a solve or an eigendecomposition of  $K$ , which takes  $O(n^3)$  time.
- Usually, we run out of memory before we run out of time.
- The practical limit on today's workstations is (more-or-less) 10,000 points (using Matlab).
- How can we do more?



# The Linear Case

- The linear kernel is  $k(X_i, X_j) = X_i^t X_j$ .
- The linear kernel offers many advantages for computation, which we now explore.
- Key idea: we get a decomposition of the kernel matrix for free:  $K = XX^t$ .
- In the linear case, we will see that we have two different computation options.

# Linear kernel, linear function

With a linear kernel, the function we are learning is linear as well:

$$\begin{aligned} f(x) &= c^t k(X, x) \\ &= c^t Xx \\ &= w^t x, \end{aligned}$$

where we define the hyperplane  $w$  to be  $X^t c$ . We can classify new points in  $O(d)$  time, using  $w$ , rather than having to compute a weighted sum of  $n$  kernel products (which will usually cost  $O(nd)$  time).

# Linear kernel, SVD approach, I

- Assume  $n$ , the number of points, is bigger than  $d$ , the number of dimensions. (If not, the best bet is to ignore the special properties of the linear kernel.)
- The economy-size SVD of  $X$  can be written as  $X = USV^t$ , with  $U \in \mathbb{R}^{n \times d}$ ,  $S \in \mathbb{R}^{d \times d}$ ,  $V \in \mathbb{R}^{d \times d}$ ,  $U^t U = V^t V = VV^t = I_d$ , and  $S$  diagonal and positive semidefinite. (Note that  $UU^t \neq I_n$ ).
- We will express the LOO formula directly in terms of the SVD, rather than  $K$ .

# Linear kernel, SVD approach, II

$$\begin{aligned}K &= XX^t = (USV^t)(VSU^t) = US^2U^t \\K + \lambda I &= US^2U^t + \lambda I_n \\&= \begin{bmatrix} U & U_{\perp} \end{bmatrix} \begin{bmatrix} S^2 + \lambda I_d & \\ & \lambda I_{n-d} \end{bmatrix} \begin{bmatrix} U^t \\ U_{\perp}^t \end{bmatrix} \\&= U(S^2 + \lambda I_d)U^t + \lambda U_{\perp}U_{\perp}^t \\&= U(S^2 + \lambda I_d)U^t + \lambda(I_n - UU^t) \\&= US^2U^t + \lambda I_n\end{aligned}$$

# Linear kernel, SVD approach, III

$$\begin{aligned}& (K + \lambda I)^{-1} \\&= (US^2U^t + \lambda I_n)^{-1} \\&= \left( \begin{bmatrix} U & U_{\perp} \end{bmatrix} \begin{bmatrix} S^2 + \lambda I_d & \\ & \lambda I_{n-d} \end{bmatrix} \begin{bmatrix} U^t \\ U_{\perp}^t \end{bmatrix} \right)^{-1} \\&= \begin{bmatrix} U & U_{\perp} \end{bmatrix} \begin{bmatrix} S^2 + \lambda I_d & \\ & \lambda I_{n-d} \end{bmatrix}^{-1} \begin{bmatrix} U^t \\ U_{\perp}^t \end{bmatrix} \\&= U(S^2 + \lambda I)^{-1}U^t + \lambda^{-1}U_{\perp}U_{\perp}^t \\&= U(S^2 + \lambda I)^{-1}U^t + \lambda^{-1}(I - UU^t) \\&= U \left[ (S^2 + \lambda I)^{-1} - \lambda^{-1}I \right] U^t + \lambda^{-1}I\end{aligned}$$

# Linear kernel, SVD approach, IV

$$\begin{aligned}c &= (K + \lambda I)^{-1} Y \\&= U \left[ (S^2 + \lambda I)^{-1} - \lambda^{-1} I \right] U^t Y + \lambda^{-1} Y\end{aligned}$$

$$G_{ij}^{-1} = \sum_{k=1}^d U_{ik} U_{jk} [(S_{kk} + \lambda)^{-1} - \lambda^{-1}] + [i = j] \lambda^{-1}$$

$$G_{ii}^{-1} = \sum_{k=1}^d U_{ik}^2 [(S_{kk} + \lambda)^{-1} - \lambda^{-1}] + \lambda^{-1}$$

$$\begin{aligned}LOOE &= \frac{c}{\text{diag}_v(G^{-1})} \\&= \frac{U \left[ (S^2 + \lambda I)^{-1} - \lambda^{-1} I \right] U^t Y + \lambda^{-1} Y}{\text{diag}_v(U \left[ (S^2 + \lambda I)^{-1} - \lambda^{-1} I \right] U^t + \lambda^{-1} I)}\end{aligned}$$

# Linear kernel, SVD approach, computational costs

- We need  $O(nd)$  memory to store the data in the first place. The (economy-sized) SVD also requires  $O(nd)$  memory, and  $O(nd^2)$  time.
- Once we have the SVD, we can compute the LOO error (for a given  $\lambda$ ) in  $O(nd)$  time.
- Compared to the nonlinear case, we have replaced an  $O(n)$  with an  $O(d)$ , in both time and memory. If  $n \gg d$ , this can represent a huge savings.

# Linear kernel, direct approach, I

For the linear kernel,

$$\begin{aligned} L &= \min_{c \in \mathbb{R}^n} \frac{1}{2} \|Y - Kc\|_2^2 + \frac{\lambda}{2} c^t Kc \\ &= \min_{c \in \mathbb{R}^n} \frac{1}{2} \|Y - XX^t c\|_2^2 + \frac{\lambda}{2} c^t XX^t c \\ &= \min_{w \in \mathbb{R}^d} \frac{1}{2} \|Y - Xw\|_2^2 + \frac{\lambda}{2} \|w\|_2^2. \end{aligned}$$

Taking the derivative with respect to  $w$ ,

$$\frac{\partial L}{\partial w} = X^t Xw - X^t Y + \lambda w,$$

and setting to zero implies

$$w = (X^t X + \lambda I)^{-1} X^t Y.$$



# Linear kernel, direct approach, II

- If we are willing to give up LOO validation, we can skip the computation of  $c$  and just get  $w$  directly.
- We can work with the *Gram matrix*  $X^t X \in \mathbb{R}^{d \times d}$ .
- The algorithm is identical to solving a general RLS problem with kernel matrix  $X^t X$  and labels  $X^t y$ .
- Form the eigendecomposition of  $X^t X$ , in  $O(d^3)$  time, form  $w(\lambda)$  in  $O(d^2)$  time.
- Why would we give up LOO validation? Maybe  $n$  is very large, so using a development set is good enough.

# Comparing the direct and SVD approaches

- Asymptotic complexity is actually the same: it takes  $O(nd^2)$  time to form the SVD of  $X$ , or to form  $X^tX$ .
- The constant in forming the SVD is about 25.
- Forming  $X^tX$  can be (relatively) easily parallelized.
- Recommendation: Use the SVD when possible, switch to the direct approach when it gets too slow.

# Introducing the Subset of Regressors

- Suppose that  $n$  is too large to apply nonlinear RLS, but we need a nonlinear kernel.
- (In some circumstances, we can *explicitly* construct nonlinear feature features, such as 2nd-order polynomial, and then use the linear approach. See my ICASSP 2007 paper.)
- Another idea is the *subset of regressors* approach.

# Subset of Regressors

- The representer theorem guarantees that the Tikhonov solution can be written as

$$f(\cdot) = \sum_{i=1}^n c_i k(X_i, \cdot),$$

for some  $c \in \mathbb{R}^n$ .

- Suppose we divide our data into two pieces,  $X_R$  and  $X_S$ , and *require a priori* that only the points in  $X_R$  have nonzero coefficients in the expansion:

$$f(\cdot) = \sum_{i=1}^{|R|} c_i k(X_i, \cdot),$$

for some  $c \in \mathbb{R}^{|R|}$ : this is the subset of regressors method.



# Subset of Regressors, Derivation

- Defining  $T = R \cup S$ , we want to find

$$\min_{c \in \mathbb{R}^n} \frac{1}{2} \|Y - K_{TR}c\|_2^2 + \frac{\lambda}{2} c^t K_{RR}c$$

.

- Setting the derivative to zero,

$$\begin{aligned} -K_{RT}Y + K_{TR}^t K_{TR}c + \lambda K_{RR}c &= 0 \\ (K_{RT}K_{TR} + \lambda K_{RR})c &= K_{RT}Y. \end{aligned}$$

# Finding $c(\lambda)$ is still cheap

Using the Cholesky factorization  $K_{RR} = GG^t$ ,

$$\begin{aligned} & K_{RT}K_{TR} + \lambda K_{RR} \\ = & K_{RT}K_{TR} + \lambda GG^t \\ = & GG^{-1}(K_{RT}K_{TR} + \lambda GG^t)G^{-t}G^t \\ = & G(G^{-1}K_{RT}K_{TR}G^{-t} + \lambda I)G^t. \end{aligned}$$

We handle varying  $\lambda$  using an eigendecomposition of  $G^{-1}K_{RT}K_{TR}G^{-t}$ .

Can we do LOO this way? Good question ...

“You should be asking how the answers will be used and what is *really* needed from the computation. Time and time again someone will ask for the inverse of a matrix when all that is needed is the solution of a linear system; for an interpolating polynomial when all that is needed is its values at some point; for the solution of an ODE at a sequence of points when all that is needed is the limiting, steady-state value. A common complaint is that least squares curve-fitting couldn't possibly work on *this* data set and some more complicated method is needed; in almost all such cases, least squares curve-fitting will work just fine because it is so very robust.”

Leader, Numerical Analysis and Scientific Computation